

# S10 T01 Duplicate

May 12, 2022

## 1 Nivell 1

### 1.1 Exercici 1

Crea almenys tres models de classificació diferents per intentar predir el millor possible l'endarreriment dels vols (ArrDelay) de DelayedFlights.csv. Considera si el vol ha arribat tard o no (ArrDelay > 0).

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: fly = pd.read_csv('C:/Users/Guillermo/Desktop/Curs Data Scientist/Sprint 2/S02_
↳T05 Exploració de les dades/DelayedFlights.csv')
```

```
[4]: fly.columns
```

```
[4]: Index(['Unnamed: 0', 'Year', 'Month', 'DayofMonth', 'DayOfWeek', 'DepTime',
'CRSDepTime', 'ArrTime', 'CRSArrTime', 'UniqueCarrier', 'FlightNum',
'TailNum', 'ActualElapsedTime', 'CRSElapsedTime', 'AirTime', 'ArrDelay',
'DepDelay', 'Origin', 'Dest', 'Distance', 'TaxiIn', 'TaxiOut',
'Cancelled', 'CancellationCode', 'Diverted', 'CarrierDelay',
'WeatherDelay', 'NASDelay', 'SecurityDelay', 'LateAircraftDelay'],
dtype='object')
```

```
[3]: fly['retard'] = fly['ArrDelay'].apply(lambda x: 1 if x > 0 else 0)
```

```
[7]: fly.head()
```

```
[7]:   Unnamed: 0  Year  Month  DayofMonth  DayOfWeek  DepTime  CRSDepTime  \
0           0  2008     1           3           4    2003.0         1955
1           1  2008     1           3           4     754.0          735
2           2  2008     1           3           4     628.0          620
3           4  2008     1           3           4    1829.0         1755
4           5  2008     1           3           4    1940.0         1915
```

```
ArrTime  CRSArrTime  UniqueCarrier  ...  TaxiOut  Cancelled  \
```

0	2211.0	2225	WN ...	8.0	0
1	1002.0	1000	WN ...	10.0	0
2	804.0	750	WN ...	17.0	0
3	1959.0	1925	WN ...	10.0	0
4	2121.0	2110	WN ...	10.0	0

	CancellationCode	Diverted	CarrierDelay	WeatherDelay	NASDelay	\
0	N	0	NaN	NaN	NaN	
1	N	0	NaN	NaN	NaN	
2	N	0	NaN	NaN	NaN	
3	N	0	2.0	0.0	0.0	
4	N	0	NaN	NaN	NaN	

	SecurityDelay	LateAircraftDelay	retard
0	NaN	NaN	0
1	NaN	NaN	1
2	NaN	NaN	1
3	0.0	32.0	1
4	NaN	NaN	1

[5 rows x 31 columns]

```
[5]: fly['retard'].isna().sum()
```

```
[5]: 0
```

```
[6]: fly['ArrDelay'].isna().sum()
```

```
[6]: 8387
```

```
[4]: #Eliminamos NaNs de ArrDelay
fly = fly.dropna(subset = ['ArrDelay'])
```

```
[5]: fly['ArrDelay'].isna().sum()
```

```
[5]: 0
```

```
[5]: fly2 = fly.sample(n = 100000, random_state = 42)
```

```
[7]: fly2.isna().sum()
```

```
[7]: Unnamed: 0      0
      Year          0
      Month         0
      DayOfMonth    0
      DayOfWeek     0
      DepTime       0
```

```

CRSDepTime      0
ArrTime         0
CRSArrTime      0
UniqueCarrier   0
FlightNum       0
TailNum         0
ActualElapsedTime 0
CRSElapsedTime  0
AirTime         0
ArrDelay        0
DepDelay        0
Origin          0
Dest            0
Distance        0
TaxiIn          0
TaxiOut         0
Cancelled       0
CancellationCode 0
Diverted        0
CarrierDelay     35303
WeatherDelay     35303
NASDelay         35303
SecurityDelay    35303
LateAircraftDelay 35303
retard           0
dtype: int64

```

Com retard no ha heretat els NaNs de Arrdelay eliminem les files amb NaNs de la nostra mostra

```
[8]: fly2.retard.value_counts()
```

```

[8]: 1    89517
     0    10483
     Name: retard, dtype: int64

```

Igual que en l'exercici anterior eliminem els NaNs de CarrierDelay, WeatherDelay, NASDelay, SecurityDelay i LateAircraftDelay amb la mitjana.

```

[6]: numVar = ['Month', 'DayofMonth', 'DayOfWeek', 'DepTime', 'CRSDepTime',
              ↪ 'ArrTime',
              'CRSArrTime', 'FlightNum', 'ActualElapsedTime', 'CRSElapsedTime',
              ↪ 'AirTime', 'ArrDelay',
              'DepDelay', 'Distance', 'TaxiIn', 'TaxiOut', 'CarrierDelay',
              'WeatherDelay', 'NASDelay', 'SecurityDelay', 'LateAircraftDelay']
catVar = ['UniqueCarrier', 'TailNum', 'Origin', 'Dest']

flyNum = fly2.loc[:, numVar]

```

```
flyCat = fly2.loc[:,catVar]
retards = fly2.loc[:, 'retard']
```

```
[10]: flyNum.head()
```

```
[10]:
```

	Month	DayofMonth	DayOfWeek	DepTime	CRSDepTime	ArrTime	\
403778	3	25	2	1124.0	1115	1221.0	
1397595	8	9	6	1152.0	1130	1555.0	
1849546	12	13	6	1226.0	1210	1405.0	
18418	1	21	1	2046.0	1950	2158.0	
1805454	12	26	5	1759.0	1709	2026.0	

	CRSArrTime	FlightNum	ActualElapsedTime	CRSElapsedTime	...	\
403778	1215	1095	57.0	60.0	...	
1397595	1520	1690	183.0	170.0	...	
1849546	1355	5615	99.0	105.0	...	
18418	2045	1274	132.0	115.0	...	
1805454	1937	6180	87.0	88.0	...	

	ArrDelay	DepDelay	Distance	TaxiIn	TaxiOut	CarrierDelay	\
403778	6.0	9.0	251	5.0	7.0	NaN	
1397595	35.0	22.0	1121	7.0	23.0	22.0	
1849546	10.0	16.0	528	7.0	13.0	NaN	
18418	73.0	56.0	629	4.0	25.0	7.0	
1805454	49.0	50.0	495	3.0	13.0	0.0	

	WeatherDelay	NASDelay	SecurityDelay	LateAircraftDelay
403778	NaN	NaN	NaN	NaN
1397595	0.0	13.0	0.0	0.0
1849546	NaN	NaN	NaN	NaN
18418	0.0	17.0	0.0	49.0
1805454	0.0	0.0	0.0	49.0

[5 rows x 21 columns]

```
[11]: flyCat.head()
```

```
[11]:
```

	UniqueCarrier	TailNum	Origin	Dest
403778	WN	N909WN	STL	MDW
1397595	AA	N5ELAA	DFW	MIA
1849546	EV	N752EV	CAK	ATL
18418	WN	N642WN	DEN	LAS
1805454	OO	N742SK	DEN	OKC

```
[12]: print(flyNum.shape)
print(flyCat.shape)
```

```
(100000, 21)
(100000, 4)
```

```
[7]: from sklearn.impute import SimpleImputer

imp = SimpleImputer(missing_values = np.nan, strategy = 'median')

flyTemp = imp.fit_transform(flyNum)

flyNum = pd.DataFrame(flyTemp, columns = numVar)

print(flyNum.isna().sum())

flyNum.head()
```

```
Month          0
DayofMonth     0
DayOfWeek      0
DepTime        0
CRSDepTime     0
ArrTime        0
CRSArrTime     0
FlightNum      0
ActualElapsedTime 0
CRSElapsedTime 0
AirTime        0
ArrDelay       0
DepDelay       0
Distance       0
TaxiIn         0
TaxiOut        0
CarrierDelay   0
WeatherDelay   0
NASDelay       0
SecurityDelay  0
LateAircraftDelay 0
dtype: int64
```

```
[7]:   Month  DayofMonth  DayOfWeek  DepTime  CRSDepTime  ArrTime  CRSArrTime  \
0     3.0         25.0         2.0   1124.0     1115.0   1221.0     1215.0
1     8.0          9.0         6.0   1152.0     1130.0   1555.0     1520.0
2    12.0         13.0         6.0   1226.0     1210.0   1405.0     1355.0
3     1.0         21.0         1.0   2046.0     1950.0   2158.0     2045.0
4    12.0         26.0         5.0   1759.0     1709.0   2026.0     1937.0

   FlightNum  ActualElapsedTime  CRSElapsedTime  ...  ArrDelay  DepDelay  \
0     1095.0             57.0             60.0  ...        6.0         9.0
```

1	1690.0	183.0	170.0	...	35.0	22.0
2	5615.0	99.0	105.0	...	10.0	16.0
3	1274.0	132.0	115.0	...	73.0	56.0
4	6180.0	87.0	88.0	...	49.0	50.0

	Distance	TaxiIn	TaxiOut	CarrierDelay	WeatherDelay	NASDelay	\
0	251.0	5.0	7.0	2.0	0.0	2.0	
1	1121.0	7.0	23.0	22.0	0.0	13.0	
2	528.0	7.0	13.0	2.0	0.0	2.0	
3	629.0	4.0	25.0	7.0	0.0	17.0	
4	495.0	3.0	13.0	0.0	0.0	0.0	

	SecurityDelay	LateAircraftDelay
0	0.0	8.0
1	0.0	0.0
2	0.0	8.0
3	0.0	49.0
4	0.0	49.0

[5 rows x 21 columns]

Igual que en el cas anterior apliquem RobustScaler i normalitzem

```
[8]: import sklearn.preprocessing as sklp

scaler = sklp.RobustScaler()

fly2NTemp = scaler.fit_transform(flyNum)

minmax = sklp.MinMaxScaler()

fly2NTemp = minmax.fit_transform(fly2NTemp)

flyNum2 = pd.DataFrame(data = fly2NTemp, columns = numVar)

print(flyNum2.shape)
flyNum2.head()
```

(100000, 21)

```
[8]:      Month  DayOfMonth  DayOfWeek  DepTime  CRSDepTime  ArrTime  \
0  0.181818    0.800000    0.166667  0.468112    0.472434  0.508545
1  0.636364    0.266667    0.833333  0.479783    0.478796  0.647770
2  1.000000    0.400000    0.833333  0.510629    0.512723  0.585244
3  0.000000    0.666667    0.000000  0.852439    0.826548  0.899125
4  1.000000    0.833333    0.666667  0.732805    0.724343  0.844102
```

	CRSArrTime	FlightNum	ActualElapsedTime	CRSElapsedTime	...	ArrDelay	\
0	0.514843	0.112332	0.064417	0.066874	...	0.038089	
1	0.644190	0.173426	0.257669	0.237947	...	0.058172	
2	0.574215	0.576445	0.128834	0.136858	...	0.040859	
3	0.866836	0.130712	0.179448	0.152411	...	0.084488	
4	0.821035	0.634459	0.110429	0.110420	...	0.067867	

	DepDelay	Distance	TaxiIn	TaxiOut	CarrierDelay	WeatherDelay	\
0	0.001980	0.044616	0.028902	0.018919	0.001434	0.0	
1	0.010561	0.221050	0.040462	0.062162	0.015771	0.0	
2	0.006601	0.100791	0.040462	0.035135	0.001434	0.0	
3	0.033003	0.121274	0.023121	0.067568	0.005018	0.0	
4	0.029043	0.094099	0.017341	0.035135	0.000000	0.0	

	NASDelay	SecurityDelay	LateAircraftDelay
0	0.001552	0.0	0.013769
1	0.010085	0.0	0.000000
2	0.001552	0.0	0.013769
3	0.013189	0.0	0.084337
4	0.000000	0.0	0.084337

[5 rows x 21 columns]

Ara apliquem dummies a les variables categoriques

```
[9]: flyCat2 = pd.get_dummies(data = flyCat, prefix = ['UC-', 'TN-', 'O-', 'D-'])
print(flyCat2.shape)
```

(100000, 5812)

```
[30]: flyCat2.head()
```

```
[30]:
```

	UC-_9E	UC-_AA	UC-_AQ	UC-_AS	UC-_B6	UC-_CO	UC-_DL	UC-_EV	\
1782417	0	0	0	0	0	0	0	0	
512712	0	0	0	0	0	0	0	0	
447137	0	0	0	0	0	0	0	0	
55082	0	0	0	0	0	0	0	0	
877634	0	0	0	0	0	1	0	0	

	UC-_F9	UC-_FL	...	D-_TXK	D-_TYR	D-_TYS	D-_VLD	D-_VPS	D-_WRG	\
1782417	0	0	...	0	0	0	0	0	0	
512712	0	0	...	0	0	0	0	0	0	
447137	0	0	...	0	0	0	0	0	0	
55082	0	0	...	0	0	0	0	0	0	
877634	0	0	...	0	0	0	0	0	0	

	D-_XNA	D-_YAK	D-_YKM	D-_YUM
1782417	0	0	0	0
512712	0	0	0	0
447137	0	0	0	0
55082	0	0	0	0
877634	0	0	0	0

[5 rows x 5811 columns]

```
[10]: #flyCat2.shape
a = flyCat2.reset_index(drop=True)
```

```
[17]: a.head()
```

```
[17]:   UC-_9E  UC-_AA  UC-_AQ  UC-_AS  UC-_B6  UC-_C0  UC-_DL  UC-_EV  UC-_F9  \
0      0      0      0      0      0      0      0      0      0
1      0      1      0      0      0      0      0      0      0
2      0      0      0      0      0      0      0      1      0
3      0      0      0      0      0      0      0      0      0
4      0      0      0      0      0      0      0      0      0
```

	UC-_FL	...	D-_TYR	D-_TYS	D-_VLD	D-_VPS	D-_WRG	D-_WYS	D-_XNA	\
0	0	...	0	0	0	0	0	0	0	
1	0	...	0	0	0	0	0	0	0	
2	0	...	0	0	0	0	0	0	0	
3	0	...	0	0	0	0	0	0	0	
4	0	...	0	0	0	0	0	0	0	

	D-_YAK	D-_YKM	D-_YUM
0	0	0	0
1	0	0	0
2	0	0	0
3	0	0	0
4	0	0	0

[5 rows x 5812 columns]

```
[18]: a.shape
```

```
[18]: (100000, 5812)
```

```
[11]: df = pd.merge(flyNum2, a, left_index=True, right_index=True)
#df = pd.concat([flyNum2, a], axis=1)
```

```
[20]: df.shape
```

```
[20]: (100000, 5833)
```



```
[21]: df.head()
```

```
[21]:      Month  DayOfMonth  DayOfWeek  DepTime  CRSDepTime  ArrTime  \
0  0.181818    0.800000    0.166667  0.468112    0.472434  0.508545
1  0.636364    0.266667    0.833333  0.479783    0.478796  0.647770
2  1.000000    0.400000    0.833333  0.510629    0.512723  0.585244
3  0.000000    0.666667    0.000000  0.852439    0.826548  0.899125
4  1.000000    0.833333    0.666667  0.732805    0.724343  0.844102

      CRSArrTime  FlightNum  ActualElapsedTime  CRSElapsedTime  ...  D-_TYR  \
0    0.514843    0.112332          0.064417          0.066874  ...      0
1    0.644190    0.173426          0.257669          0.237947  ...      0
2    0.574215    0.576445          0.128834          0.136858  ...      0
3    0.866836    0.130712          0.179448          0.152411  ...      0
4    0.821035    0.634459          0.110429          0.110420  ...      0

      D-_TYS  D-_VLD  D-_VPS  D-_WRG  D-_WYS  D-_XNA  D-_YAK  D-_YKM  D-_YUM
0         0         0         0         0         0         0         0         0         0
1         0         0         0         0         0         0         0         0         0
2         0         0         0         0         0         0         0         0         0
3         0         0         0         0         0         0         0         0         0
4         0         0         0         0         0         0         0         0         0
```

[5 rows x 5833 columns]

```
[22]: print(flyCat2.shape)
print(flyNum2.shape)
print(df.shape)
```

(100000, 5812)

(100000, 21)

(100000, 5833)

Com tenim un dataset amb un total de 5833 columnes passem a reduir-lo aplicant una PCA

### 1.1.1 PCA

```
[12]: retards = retards.reset_index(drop=True)
retards
```

```
[12]: 0      1
1      1
2      1
3      1
4      1
...
99995  1
99996  1
```

```
99997    1
99998    1
99999    0
Name: retard, Length: 100000, dtype: int64
```

```
[25]: retards.value_counts()
```

```
[25]: 1    89517
      0    10483
      Name: retard, dtype: int64
```

```
[28]: from sklearn.decomposition import PCA

      #pca = PCA(.95)#Retenga el 95% de la informaci3n
      pca = PCA().fit(df)
```

```
-----
MemoryError                                Traceback (most recent call last)
C:\Users\GUILLER~1\AppData\Local\Temp\ipykernel_7920\1192173359.py in <module>
      3 #pca = PCA(.95)#Retenga el 95% de la informaci3n
      4
----> 5 pca = PCA().fit(df)

~\anaconda3\lib\site-packages\sklearn\decomposition\_pca.py in fit(self, X, y)
    380         Returns the instance itself.
    381         """
--> 382         self._fit(X)
    383         return self
    384

~\anaconda3\lib\site-packages\sklearn\decomposition\_pca.py in _fit(self, X)
    455         # Call different fits for either full or truncated SVD
    456         if self._fit_svd_solver == "full":
--> 457             return self._fit_full(X, n_components)
    458         elif self._fit_svd_solver in ["arpack", "randomized"]:
    459             return self._fit_truncated(X, n_components, self.
-> _fit_svd_solver)

~\anaconda3\lib\site-packages\sklearn\decomposition\_pca.py in _fit_full(self, X, n_components)
    490         X -= self.mean_
    491
--> 492         U, S, Vt = linalg.svd(X, full_matrices=False)
    493         # flip eigenvectors' sign to enforce deterministic output
    494         U, Vt = svd_flip(U, Vt)
```

```

~\anaconda3\lib\site-packages\scipy\linalg\decomp_svd.py in svd(a,
↳ full_matrices, compute_uv, overwrite_a, check_finite, lapack_driver)
    125
    126     # perform decomposition
--> 127     u, s, v, info = gesXd(a1, compute_uv=compute_uv, lwork=lwork,
    128                           full_matrices=full_matrices,
↳ overwrite_a=overwrite_a)
    129

MemoryError: Unable to allocate 4.35 GiB for an array with shape (100000, 5833),
↳ and data type float64

```

Reduïm la mostra per intentar processar la PCA

```
[13]: df2 = pd.merge(df, retards, left_index=True, right_index=True)
df2.shape
```

```
[13]: (100000, 5834)
```

```
[14]: dfReduct = df2.sample(n=10000, random_state=42)
dfReduct.shape
```

```
[14]: (10000, 5834)
```

```
[16]: from sklearn.decomposition import PCA
pca = PCA(.95)
retardsReduct = dfReduct.loc[:, 'retard']
```


```
[17]: dfReduct = dfReduct.drop('retard', axis=1)
```

```
[18]: pca.fit(dfReduct)
dfPCA = pca.transform(dfReduct)
```

```
[19]: dfPCA.shape
```


```
[19]: (10000, 2664)
```

```
[20]: sns.set()
plt.plot(np.cumsum(pca.explained_variance_ratio_))
plt.xlabel('número de components')
plt.ylabel('cumulative explained variance')
plt.show()
```



Un cop reduït el número de variables mirem com de desbalancejats estan els models

```
[39]: retardsReduct.value_counts().plot(kind='bar')  
plt.show()
```



Com el model està molt desbalancejat convé equilibrar-lo (Cal dividir primer el dataset i després rebalancejar).

```
[21]: from sklearn.model_selection import train_test_split

x_train, x_test, y_train, y_test = train_test_split(dfPCA, retardsReduct,
    ↳test_size = 0.25, random_state = 42)

print(x_train.shape, '\n', y_train.shape, '\n', x_test.shape, '\n', y_test.
    ↳shape)
```

```
(7500, 2664)
(7500,)
(2500, 2664)
(2500,)
```

Com el dataset és molt gran una opció seria agafar més mostra. Però degut als problemes de memòria serà una millor opció generar dades noves a partir de la PCA amb SMOUT

```
[43]: y_train.value_counts()
```

```
[43]: 1    6691
      0     809
      Name: retard, dtype: int64
```

```
[22]: from imblearn.over_sampling import SMOTE

x_resampled, y_resampled = SMOTE().fit_resample(x_train, y_train)
```

```
[23]: y_resampled.value_counts()
```

```
[23]: 1    6691
      0    6691
      Name: retard, dtype: int64
```

```
[46]: print(imblearn.__version__)# No sé si m'he de preocupar que això funcioni o no
```

```
-----
NameError                                Traceback (most recent call last)
C:\Users\GUILLE~1\AppData\Local\Temp\ipykernel_7920\2763361883.py in <module>
----> 1 print(imblearn.__version__)# No sé si m'he de preocupar que això
    ↳funcioni o no

NameError: name 'imblearn' is not defined
```

Amb tots els processos realitzats hem escalat les variables, reduït les dimensions del dataframe i en

el cas del test que hi hagi la mateixa proporció de vols amb retard i sense. Ara passem a aplicar els models.

### 1.1.2 Models a aplicar


- Decision tree
- Suport vector machine
- Logistic regression

```
[25]: # Decision tree
from sklearn.tree import DecisionTreeClassifier

dtc = DecisionTreeClassifier()
tree = dtc.fit(x_resampled, y_resampled)
tree_predict = tree.predict(x_test)
```

```
[33]: #SVM

sns.scatterplot(x=dfPCA[:,0], y=dfPCA[:,1], hue=retardsReduct )
plt.xlabel('1r component')
plt.ylabel('2n component')
plt.title('PCA 1r i 2n component')
plt.show()
```



La gràfica mostra que hi han dos grups diferenciats però no sembla que amb els dos primers

components sigui degut els grups pels retards. Igualment s'intentarà separar-los per SVG amb kernel lineal

```
[34]: from sklearn import svm
```

```
clf = svm.SVC(kernel='linear')
clf.fit(x_resampled, y_resampled)
svm_predict = clf.predict(x_test)
```

```
[35]: #Logistic regression
```

```
from sklearn.linear_model import LogisticRegression
```

```
logreg = LogisticRegression()
```

```
logreg.fit(x_resampled, y_resampled)
```

```
logreg_predict = logreg.predict(x_test)
```

```
C:\Users\Guillermo\anaconda3\lib\site-
packages\sklearn\linear_model\_logistic.py:814: ConvergenceWarning: lbfgs failed
to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

```
[39]: print(logreg_predict.shape)
print(x_test.shape)
```

```
(2500,)
```

```
(2500, 2664)
```

Ignorarem l'alerta degut a que a les dades ja hem aplicat un robust scaler i han sigut normalitzades. Tot i així, al no poder fer prou iteracions pot donar un resultat dolent

## 1.2 Exercici 2

Compara els models de classificació utilitzant la precisió (accuracy), una matriu de confiança i d'altres mètriques més avançades.

```
[44]: from sklearn import metrics
```

```
#Accuracy
```

```
print('decision tree', metrics.accuracy_score(y_test, tree_predict))
```

```
print('SVM', metrics.accuracy_score(y_test, svm_predict))
```

```
print('Logistic regression', metrics.accuracy_score(y_test, logreg_predict))
```

```
decision tree 0.842
SVM 0.8396
Logistic regression 0.824
```

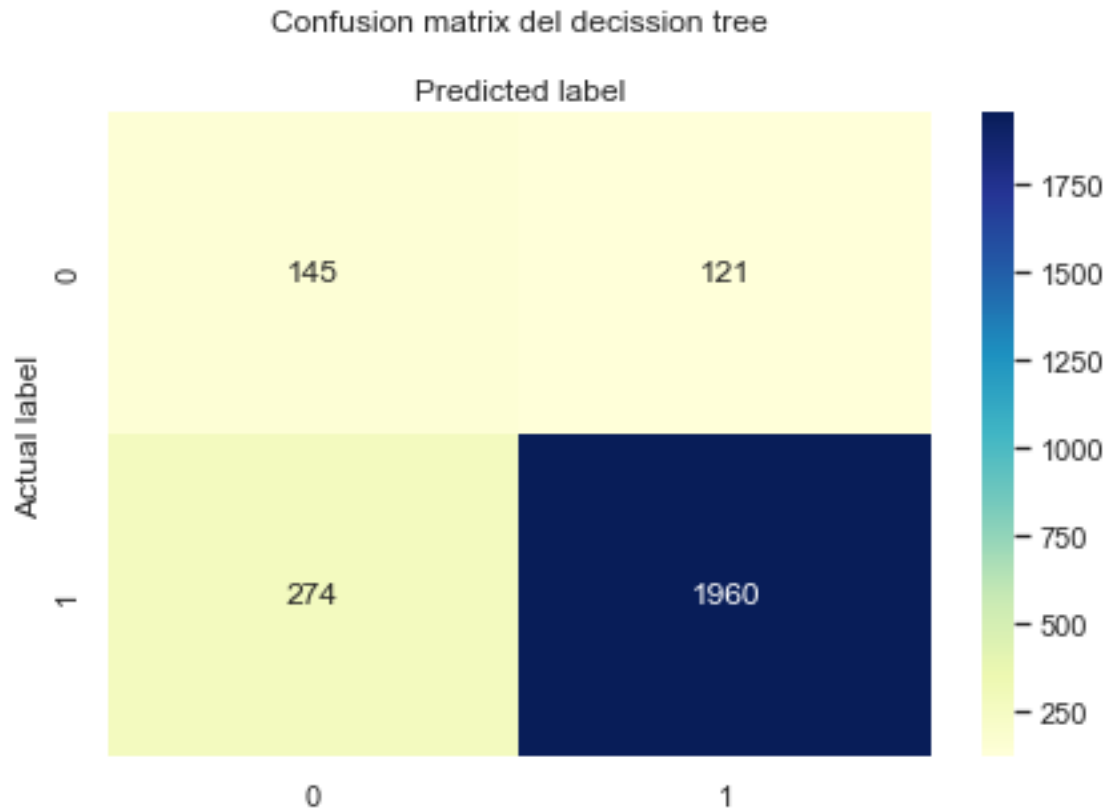
### 1.2.1 Matriu de confiança

```
[47]: #Decision tree
tree_matrix = metrics.confusion_matrix(y_test, tree_predict)
tree_matrix
```

```
[47]: array([[ 145,  121],
           [ 274, 1960]], dtype=int64)
```

```
[50]: class_names=[0,1] # name of classes
fig, ax = plt.subplots()
tick_marks = np.arange(len(class_names))
plt.xticks(tick_marks, class_names)
plt.yticks(tick_marks, class_names)
# create heatmap
sns.heatmap(pd.DataFrame(tree_matrix), annot=True, cmap="YlGnBu", fmt='g')
ax.xaxis.set_label_position("top")
plt.tight_layout()
plt.title('Confusion matrix del decission tree', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
plt.show()
```






```
[53]: # SVM
svm_matrix = metrics.confusion_matrix(y_test, svm_predict)
svm_matrix
```

```
[53]: array([[ 147,  119],
        [ 282, 1952]], dtype=int64)
```


```
[54]: class_names=[0,1] # name of classes
fig, ax = plt.subplots()
tick_marks = np.arange(len(class_names))
plt.xticks(tick_marks, class_names)
plt.yticks(tick_marks, class_names)
# create heatmap
sns.heatmap(pd.DataFrame(svm_matrix), annot=True, cmap="YlGnBu", fmt='g')
ax.xaxis.set_label_position("top")
plt.tight_layout()
plt.title('Confusion matrix del svm', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
plt.show()
```



```
[57]: #Logistic regression
lgr_matrix = metrics.confusion_matrix(y_test,logreg_predict)
lgr_matrix
```

```
[57]: array([[ 149,  117],
        [ 323, 1911]], dtype=int64)
```

```
[58]: class_names=[0,1] # name of classes
fig, ax = plt.subplots()
tick_marks = np.arange(len(class_names))
plt.xticks(tick_marks, class_names)
plt.yticks(tick_marks, class_names)
# create heatmap
sns.heatmap(pd.DataFrame(lgr_matrix), annot=True, cmap="YlGnBu" ,fmt='g')
ax.xaxis.set_label_position("top")
plt.tight_layout()
plt.title('Confusion matrix del logistic regression', y=1.1)
plt.ylabel('Actual label')
plt.xlabel('Predicted label')
plt.show()
```



```
[63]: data = {'models': ['tree', 'SVM', 'logreg'],
              'TP': [tree_matrix[1,1], svm_matrix[1,1], lgr_matrix[1,1]],
              'FP': [tree_matrix[0,1], svm_matrix[0,1], lgr_matrix[0,1]],
              'TN': [tree_matrix[0,0], svm_matrix[0,0], lgr_matrix[0,0]],
              'FN': [tree_matrix[1,0], svm_matrix[1,0], lgr_matrix[1,0]]
            }
resum = pd.DataFrame(data = data)
resum
```

```
[63]:
```

	models	TP	FP	TN	FN
0	tree	1960	121	145	274
1	SVM	1952	119	147	282
2	logreg	1911	117	149	323

```
[84]: def clasif_statistic(df):
        for index, row in resum.iterrows():
            print()
            precision = row[1]/(row[1]+row[2])
            sensitivity = row[1]/(row[1]+row[4])
            specificity = row[3]/(row[3]+row[2])
            print(row[0], 'Precission = ', round(precision,3))
```

```
print(row[0], 'Sensitivity = ', round(sensitivity,3))
print(row[0], 'Specificity = ', round(specificity,3))
```

```
[85]: clasif_statistic(resum)
```

```
tree Precission = 0.942
tree Sensitivity = 0.877
tree Specificity = 0.545
```

```
SVM Precission = 0.943
SVM Sensitivity = 0.874
SVM Specificity = 0.553
```

```
logreg Precission = 0.942
logreg Sensitivity = 0.855
logreg Specificity = 0.56
```

Veiem que en general la precisió i l'assignació de positius és millor que la de negatius. Sent els millors models decision tree i SVM.

### 1.3 Exercici 3

Entrena'ls utilitzant els diferents paràmetres que admeten.

```
[90]: # Decision tree
from sklearn.tree import DecisionTreeClassifier

dtc2 = DecisionTreeClassifier(criterion = 'entropy')
tree2 = dtc2.fit(x_resampled, y_resampled)
tree_predict2 = tree2.predict(x_test)

metrics.accuracy_score(y_test, tree_predict2)
matrix = metrics.confusion_matrix(y_test, tree_predict2)
```

```
[94]: def clasif_statistic2(m):
    precision = m[1,1]/(m[1,1]+m[0,1])
    sensitivity = m[1,1]/(m[1,1]+m[1,0])
    specificity = m[0,0]/(m[0,0]+m[0,1])
    print('TP: ', m[1,1])
    print('TN: ', m[0,0])
    print('FP: ', m[0,1])
    print('FN: ', m[1,0])
    print('Precission = ', round(precision,3))
    print('Sensitivity = ', round(sensitivity,3))
```

```

        print('Specificity = ', round(specificity,3))

clasif_statistic2(matrix)
print('accuracy: ', metrics.accuracy_score(y_test, tree_predict2))

```

```

TP: 1959
TN: 163
FP: 103
FN: 275
Precision = 0.95
Sensitivity = 0.877
Specificity = 0.613
accuracy: 0.8488

```

```

[95]: clf2 = svm.SVC(kernel='rbf', gamma = 0.1)
      clf2.fit(x_resampled, y_resampled)
      svm_predict2 = clf2.predict(x_test)

      metrics.accuracy_score(y_test, svm_predict2)
      matrixSVM = metrics.confusion_matrix(y_test,svm_predict2)

      clasif_statistic2(matrixSVM)
      print('accuracy: ', metrics.accuracy_score(y_test, svm_predict2))

```

```

TP: 2109
TN: 65
FP: 201
FN: 125
Precision = 0.913
Sensitivity = 0.944
Specificity = 0.244
accuracy: 0.8696

```

```

[100]: #Logistic regression
      from sklearn.linear_model import LogisticRegression

      logreg2 = LogisticRegression(solver = 'saga')

      logreg2.fit(x_resampled, y_resampled)

      logreg_predict2 = logreg2.predict(x_test)

      matrix_lgrg = metrics.confusion_matrix(y_test, logreg_predict2)
      clasif_statistic2(matrix_lgrg)
      print('accuracy: ', metrics.accuracy_score(y_test, svm_predict2))

```

```

TP: 1911
TN: 149

```

```
FP: 117
FN: 323
Precision = 0.942
Sensitivity = 0.855
Specificity = 0.56
accuracy: 0.8696
```

## 1.4 Exercici 4

Compara el seu rendiment utilitzant l'aproximació train/test o cross-validation.

S'utilitzarà cross-validation en els models emprats en l'exercici 1

```
[102]: from sklearn.model_selection import cross_val_score

#Decision tree

scores_tree = cross_val_score(dtc, x_train, y_train, cv=5)
print(scores_tree)
print('mitja', scores_tree.mean())
print('sd', scores_tree.std())
```

```
[0.86733333 0.86533333 0.86666667 0.87266667 0.86133333]
mitja 0.8666666666666666
sd 0.003651483716701135
```

```
[103]: #SVC

scores_SVC = cross_val_score(clf, x_train, y_train, cv=5)
print(scores_SVC)
print('mitja ', scores_SVC.mean())
print('std ', scores_SVC.std())
```

```
[0.88666667 0.89133333 0.88866667 0.88133333 0.882      ]
mitja 0.8859999999999999
std 0.003841296077570124
```

```
[104]: #Logistic regression, utilitzem el cas de l'exercici 2 perquè no ha donat
#cap alerta

scores_logreg2 = cross_val_score(logreg2, x_train, y_train, cv=5)
print(scores_logreg2)
print('mean: ', scores_logreg2.mean())
print('std: ', scores_logreg2.std())
```

```
[0.89333333 0.89333333 0.894      0.89066667 0.89133333]
mean: 0.8925333333333334
std: 0.0012927146286443427
```

En tots els casos s'obtenen valors per sobre del 80% d'accuracy amb una dispersió molt petita.

[ ]: